

Lab Cloud Composer

Setting (google cloud shell, Kubernetes Engine API, cloud composer)

- ✓ Activate Google Cloud Shell
 - * gcloud auth list
 - * gcloud config list project
- ✓ Check project permissions
- ✓ Ensure that the Kubernetes Engine API is successfully enabled
Kubernetes Engine API > Manage > Disable API > Enable
- ✓ Ensure that the Cloud Composer API is successfully enabled
Cloud Composer API > Enable
- ✓ **Cloud Storage > Browser > Create bucket.**(output for the Hadoop job)
- ✓ **Navigation menu > Composer > select composer1 and set properties for environment**
click create (10-20 minutes)

Property	Value
Name	highcpu
Location	us-central1
Zone	us-central1-a
Machine type	n1-highcpu-4

Airflow and core concepts

While waiting for your Composer environment to get created, review some terms that are used with Airflow.

[Airflow](#) is a platform to programmatically author, schedule and monitor workflows.

Use Airflow to author workflows as directed acyclic graphs (DAGs) of tasks. The airflow scheduler executes your tasks on an array of workers while following the specified dependencies.

Core concepts

[DAG](#)

A Directed Acyclic Graph is a collection of all the tasks you want to run, organized in a way that reflects their relationships and dependencies.

[Operator](#)

The description of a single task, it is usually atomic. For example, the *BashOperator* is used to execute bash command.

[Task](#)

A parameterised instance of an Operator; a node in the DAG.

[Task Instance](#)

A specific run of a task; characterized as: a DAG, a Task, and a point in time. It has an indicative state: *running*, *success*, *failed*, *skipped*, ...

Defining the workflow

Now let's discuss the workflow you'll be using. Cloud Composer workflows are comprised of [DAGs \(Directed Acyclic Graphs\)](#). DAGs are defined in standard Python files that are placed in Airflow's DAG_FOLDER. Airflow will execute the code in each file to dynamically build the DAG objects. You can have as many DAGs as you want, each describing an arbitrary number of tasks. In general, each one should correspond to a single logical workflow.

Below is the `hadoop_tutorial.py` workflow code, also referred to as the DAG:

```
"""Example Airflow DAG that creates a Cloud Dataproc cluster,  
runs the Hadoop
```

```
wordcount example, and deletes the cluster.
```

```
This DAG relies on three Airflow variables
```

```
https://airflow.apache.org/concepts.html#variables
```

```
* gcp_project - Google Cloud Project to use for the Cloud  
Dataproc cluster.
```

```
* gce_zone - Google Compute Engine zone where Cloud  
Dataproc cluster should be  
created.
```

```
* gcs_bucket - Google Cloud Storage bucket to used as output  
for the Hadoop jobs from Dataproc.
```

```
See https://cloud.google.com/storage/docs/creating-buckets  
for creating a  
bucket.
```

```
"""
```

```
import datetime
```

```
import os
```

```
from airflow import models
```

```
from airflow.contrib.operators import dataproc_operator
```

```
from airflow.utils import trigger_rule
```

```
# Output file for Cloud Dataproc job.
```

```
output_file = os.path.join(  
    models.Variable.get('gcs_bucket'), 'wordcount',  
    datetime.datetime.now().strftime('%Y%m%d-%H%M%S')) +
```

```
os.sep
```

```
# Path to Hadoop wordcount example available on every  
Dataproc cluster.
```

```
WORDCOUNT_JAR = (
```

```
    'file:///usr/lib/hadoop-mapreduce/hadoop-mapreduce-  
examples.jar'
```

```
)
```

```
# Arguments to pass to Cloud Dataproc job.
wordcount_args = ['wordcount',
'gs://pub/shakespeare/rose.txt', output_file]
yesterday = datetime.datetime.combine(
    datetime.datetime.today() - datetime.timedelta(1),
    datetime.datetime.min.time())
default_dag_args = {
    # Setting start date as yesterday starts the DAG immediately
    when it is
    # detected in the Cloud Storage bucket.
    'start_date': yesterday,
    # To email on failure or retry set 'email' arg to your email and
    enable
    # emailing here.
    'email_on_failure': False,
    'email_on_retry': False,
    # If a task fails, retry it once after waiting at least 5 minutes
    'retries': 1,
    'retry_delay': datetime.timedelta(minutes=5),
    'project_id': models.Variable.get('gcp_project')
}
```

```
with models.DAG(
    'composer_sample_quickstart',
    # Continue to run DAG once per day
    schedule_interval=datetime.timedelta(days=1),
    default_args=default_dag_args) as dag:
    # Create a Cloud Dataproc cluster.
    create_dataproc_cluster =
    dataproc_operator.DataprocClusterCreateOperator(
        task_id='create_dataproc_cluster',
        # Give the cluster a unique name by appending the date
        scheduled.
        # See https://airflow.apache.org/code.html#default-variables
        cluster_name='composer-hadoop-tutorial-cluster-{{
        ds_nodash }}',
        num_workers=2,
        region='us-central1',
        zone=models.Variable.get('gce_zone'),
        image_version='2.0',
        master_machine_type='n1-standard-2',
        worker_machine_type='n1-standard-2')
```

```
# Run the Hadoop wordcount example installed on the Cloud
Dataproc cluster
  # master node.
  run_dataproc_hadoop =
dataproc_operator.DataProcHadoopOperator(
    task_id='run_dataproc_hadoop',
    region='us-central1',
    main_jar=WORDCOUNT_JAR,
    cluster_name='composer-hadoop-tutorial-cluster-{{
ds_nodash }}',
    arguments=wordcount_args)
  # Delete Cloud Dataproc cluster.
  delete_dataproc_cluster =
dataproc_operator.DataprocClusterDeleteOperator(
    task_id='delete_dataproc_cluster',
    region='us-central1',
    cluster_name='composer-hadoop-tutorial-cluster-{{
ds_nodash }}',
    # Setting trigger_rule to ALL_DONE causes the cluster to
be deleted
    # even if the Dataproc job fails.
    trigger_rule=trigger_rule.TriggerRule.ALL_DONE)
  # Define DAG dependencies.
  create_dataproc_cluster >> run_dataproc_hadoop >>
delete_dataproc_cluster
```

DAG

To orchestrate the three workflow tasks, the DAG imports the following **operators**:

1.DataprocClusterCreateOperator: Creates a Cloud Dataproc cluster.

2.DataProcHadoopOperator: Submits a Hadoop wordcount job and writes results to a Cloud Storage bucket.

3.DataprocClusterDeleteOperator: Deletes the cluster to avoid incurring ongoing Compute Engine charges.

The tasks run sequentially, which you can see in this section of the file:

Define DAG dependencies.

```
create_dataproc_cluster >> run_dataproc_hadoop >> delete_dataproc_cluster
```

The name of the DAG is quickstart, and the DAG runs once each day.

```
with models.DAG(  
    'composer_sample_quickstart',  
    # Continue to run DAG once per day  
    schedule_interval=datetime.timedelta(days=1),  
    default_args=default_dag_args) as dag:
```

Because the `start_date` that is passed in to `default_dag_args` is set to yesterday, Cloud Composer schedules the workflow to start immediately after the DAG uploads.

Viewing environment information

Go back to **Composer** to check the status of your environment.

Once your environment has been created, click the name of the environment (**highcpu**) to see its details.

On the Environment details you'll see information such as the Airflow web interface URL, Kubernetes Engine cluster ID, and a link to the DAGs folder, which is stored in your bucket.

Airflow UI ([environment variables](#))

1. Go back to the **Environments** page
2. In the **Airflow webserver** column for the environment, click **Airflow**.
3. Click on your lab credentials.
4. The Airflow web interface opens in a new browser window.

Setting Airflow variables

Airflow variables are an Airflow-specific concept that is distinct from [environment variables](#).

1. Select **Admin** > **Variables** from the Airflow menu bar, then **Create**.

The screenshot shows the Airflow web interface. At the top, a teal navigation bar contains the Airflow logo and several menu items: DAGs, Data Profiling, Browse, Admin, Docs, and About. The 'Admin' menu item is highlighted with a red rectangular box. Below the navigation bar, the main content area is titled 'Variables'. Under this title, there are two buttons: 'Choose File' (which shows 'No file chosen') and 'Import Variables'. Below these buttons, there is a horizontal bar containing a 'List' button, a 'Create' button (highlighted with a red rectangular box), an 'Add Filter' dropdown, a 'With selected' dropdown, and a search input field. Below this bar is a table with four columns: a checkbox column, a 'Key' column, a 'Val' column, and an 'Is Encrypted' column. The table is currently empty, and a message at the bottom right states 'There are no items in the table.'

Create the following Airflow variables, gcp_project, gcs_bucket, and gce_zone:

KEY	VALUE	Details
gcp_project	<your project-id>	The Google Cloud Platform project you're using for this quickstart.
gcs_bucket	gs://<my-bucket>	Replace <my-bucket> with the name of the Cloud Storage bucket you made earlier. This bucket stores the output from the Hadoop jobs from Dataproc.
gce_zone	us-central1-a	This is the Compute Engine zone where your Cloud Dataproc cluster will be created. To chose a different zone, see Available regions & zones.

Click **Save and Add Another** after adding first two variable and click **Save** for the third variable. Your Variables table should look like this when you're finished:

Variables

Choose File

No file chosen

Import Variables




List (3)

Create

Add Filter▼

With selected▼

Search

<input type="checkbox"/>		Key	Val
<input type="checkbox"/>	 	gcp_project	project-id
<input type="checkbox"/>	 	gcs_bucket	gs://bucket-name
<input type="checkbox"/>	 	gce_zone	zone

Uploading the DAG to Cloud Storage

In Cloud Shell, upload a copy of the `hadoop_tutorial.py` file to the Cloud Storage bucket that was automatically created when you created the environment.

Replace `<DAGs_folder_path>` in the following command:


```
gsutil cp gs://cloud-training/datawarehousing/lab_assets/hadoop_tutorial.py <DAGs_folder_path>
```


with the path to the DAGs folder. You can get the path by going to **Composer**. Click on the environment you created earlier and then click on the **Environment Configuration** tab to see the details of the environment. Find DAGs folder and copy the path.


```
gsutil cp gs://cloud-training/datawarehousing/lab_assets/hadoop_tutorial.py gs://us-central1-highcpu-0682d8c0-bucket/dag
```


Web server configuration


Network access control	All IP addresses have access (default) EDIT
Machine type	composer-n1-webserver-2 (2 vCPU, 1.6 GB memory) EDIT
DAGs folder	gs://us-central1-highcpu-0682d8c0-bucket/dags
Airflow web UI	https://d0ebf68b3eaa37da0p-tp.appspot.com


Storage


Browser


Monitoring

Transfer

Transfer for on-premises


Transfer Appliance

Settings


Bucket details



us-central1-highcpu-0682d8c0-bucket

OBJECTS CONFIGURATION PERMISSIONS

Buckets > us-central1-highcpu-0682d8c0-bucket > dags 

UPLOAD FILES UPLOAD FOLDER CREATE FOLDER

 Filter by object or folder name prefix


<input type="checkbox"/>	Name	Size	Type
<input type="checkbox"/>	 airflow_monitoring.py	729 B	text/x-python
<input type="checkbox"/>	 hadoop_tutorial.py	4.1 KB	text/x-python

When a DAG file is added to the DAGs folder, Cloud Composer adds the DAG to Airflow and schedules it automatically. DAG changes occur within 3-5 minutes.

You can see the task status of the composer_hadoop_tutorial DAG in the Airflow web interface.

Exploring DAG runs

When you upload your DAG file to the dags folder in Cloud Storage, Cloud Composer parses the file. If no errors are found, the name of the workflow appears in the DAG listing, and the workflow is queued to run immediately. Make sure that you're on the DAGs tab in the Airflow web interface. It takes several minutes for this process to complete. Refresh your browser to make sure you're looking at the latest information.

 Airflow

DAGs

Data Profiling ▾

Browse ▾


Admin ▾

Docs ▾

About ▾




























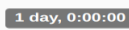



















highcpu






14:35 UTC



DAGs

Search:

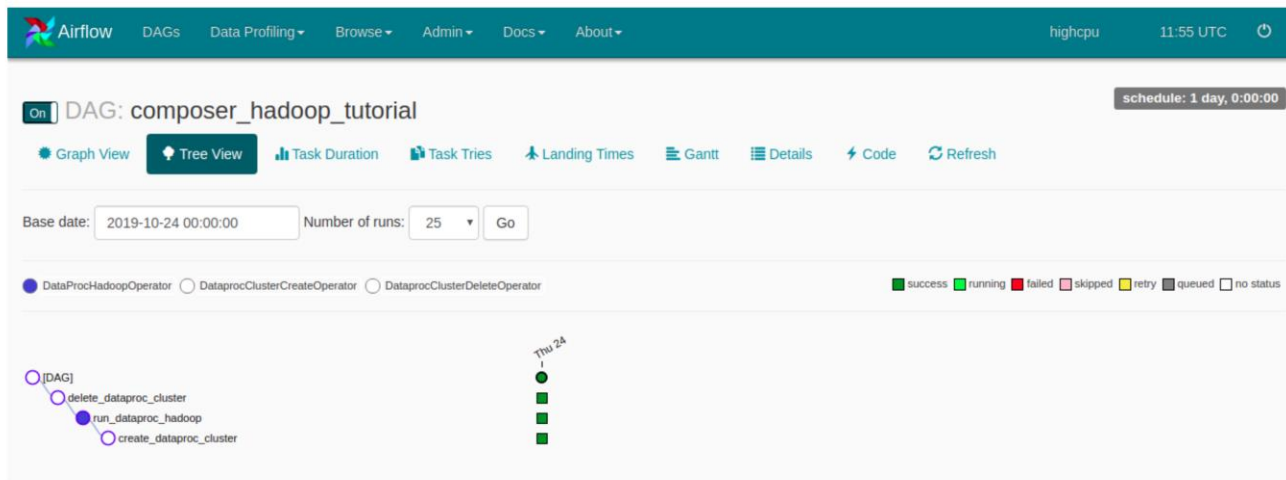
		DAG	Schedule	Owner	Recent Tasks 	Last Run 	DAG Runs 	Links
		airflow_monitoring		Airflow	    		  	         
		composer_hadoop_tutorial		Airflow	    	2019-10-24 00:00 	  	         

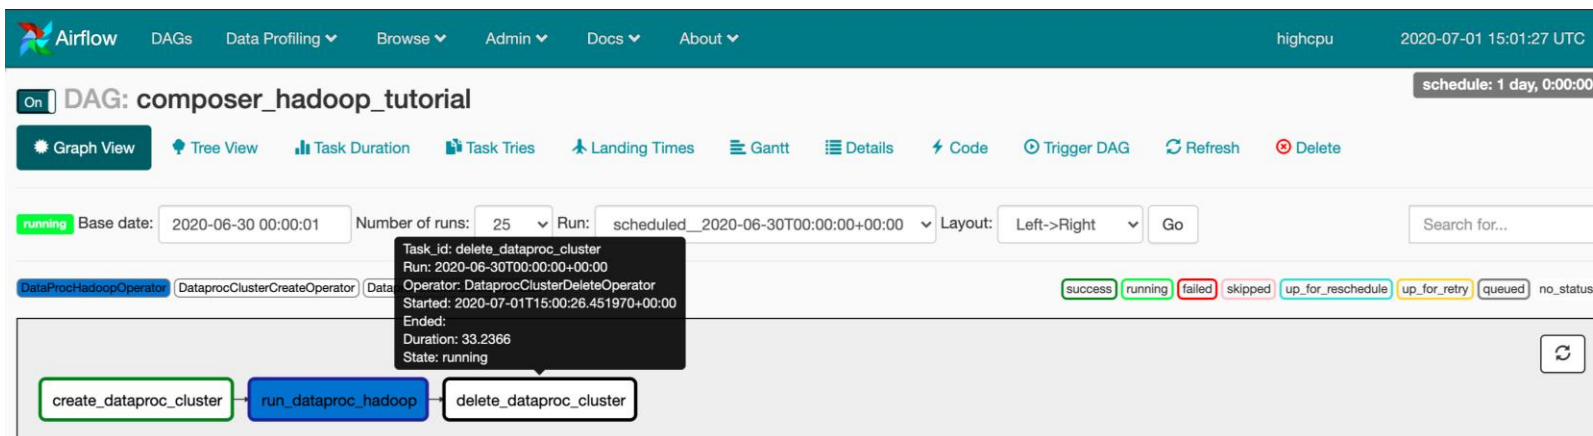
[Hide Paused DAGs](#)

Showing 1 to 2 of 2 entries

1. In Airflow, click **composer_hadoop_tutorial** to open the DAG details page. This page includes several representations of the workflow tasks and dependencies.



2. In the toolbar, click **Graph View**. Mouseover the graphic for each task to see its status. Note that the border around each task also indicates the status (green border = running; red = failed, etc.).



3. Click the "Refresh" link to make sure you're looking at the most recent information. The borders of the processes change colors as the state of the process changes

Note: If your Dataproc cluster already exists, you can run the workflow again to reach the success state by clicking `create_dataproc_cluster` graphic and then click **Clear** to reset the three tasks and click **OK** to confirm.

4. Once the status for **create_dataproc_cluster** has changed to "running", go to **Navigation menu > Dataproc**, then click on:

- **Clusters** to monitor cluster creation and deletion. The cluster created by the workflow is ephemeral: it only exists for the duration of the workflow and is deleted as part of the last workflow task.

- **Jobs** to monitor the Apache Hadoop wordcount job. Click the Job ID to see job log output.

5. Once Dataproc gets to a state of "Running", return to Airflow and click **Refresh** to see that the cluster is complete.

When the `run_dataproc_hadoop` process is complete, go to **Navigation menu > Cloud Storage > Browser** and click on the name of your bucket to see the results of the wordcount in the wordcount folder.

6. Once all the steps are complete in the DAG, each step has a dark green border. Additionally the Dataproc cluster that was created is now deleted.

```
gsutil cp gs://cloud-  
training/datawarehousing/lab_assets/hadoop_tutorial.py  
gs://us-central1-highcpu-14592db6-bucket/dags
```